

(<https://google.com/racialequity?authuser=0>)

Regularization for Simplicity: L_2 Regularization

Estimated Time: 7 minutes

Consider the following **generalization curve**, which shows the loss for both the training set and validation set against the number of training iterations.




Figure 1. Loss on training set and validation set.

Figure 1 shows a model in which training loss gradually decreases, but validation loss eventually goes up. In other words, this generalization curve shows that the model is overfitting.

(<https://developers.google.com/machine-learning/crash-course/generalization/peril-of-overfitting?authuser=0>)

to the data in the training set. Channeling our inner Ockham

(<https://developers.google.com/machine-learning/crash-course/generalization/peril-of-overfitting?authuser=0#ockham>)

, perhaps we could prevent overfitting by penalizing complex models, a principle called **regularization**.

In other words, instead of simply aiming to minimize loss (empirical risk minimization):

$$\text{minimize}(\text{Loss}(\text{Data}|\text{Model}))$$

we'll now minimize loss+complexity, which is called **structural risk minimization**:

$$\text{minimize}(\text{Loss}(\text{Data}|\text{Model}) + \text{complexity}(\text{Model}))$$

Our training optimization algorithm is now a function of two terms: the **loss term**, which measures how well the model fits the data, and the **regularization term**, which measures model complexity.

Machine Learning Crash Course focuses on two common (and somewhat related) ways to think of model complexity:

- Model complexity as a function of the *weights* of all the features in the model.
- Model complexity as a function of the *total number of features* with nonzero weights.
(A later module (<https://developers.google.com/machine-learning/crash-course/regularization-for-sparsity/l1-regularization?authuser=0>) covers this approach.)

If model complexity is a function of weights, a feature weight with a high absolute value is more complex than a feature weight with a low absolute value.

We can quantify complexity using the **L_2 regularization** formula, which defines the regularization term as the sum of the squares of all the feature weights:

$$L_2 \text{ regularization term} = ||\mathbf{w}||_2^2 = w_1^2 + w_2^2 + \dots + w_n^2$$

In this formula, weights close to zero have little effect on model complexity, while outlier weights can have a huge impact.

For example, a linear model with the following weights:

$$\{w_1 = 0.2, w_2 = 0.5, w_3 = 5, w_4 = 1, w_5 = 0.25, w_6 = 0.75\}$$

Has an L_2 regularization term of 26.915:

$$\begin{aligned} &w_1^2 + w_2^2 + \mathbf{w_3^2} + w_4^2 + w_5^2 + w_6^2 \\ &= 0.2^2 + 0.5^2 + \mathbf{5^2} + 1^2 + 0.25^2 + 0.75^2 \\ &= 0.04 + 0.25 + \mathbf{25} + 1 + 0.0625 + 0.5625 \\ &= 26.915 \end{aligned}$$

But w_3 (bolded above), with a squared value of 25, contributes nearly all the complexity. The sum of the squares of all five other weights adds just 1.915 to the L_2 regularization term.

irms

Generalization curve

https://developers.google.com/machine-learning/glossary?authuser=0#generalization_curve

Overfitting

<https://developers.google.com/machine-learning/glossary?authuser=0#overfitting>

Structural risk minimization

<https://developers.google.com/machine-learning/glossary?authuser=0#SRM>

- L_2 regularization

https://developers.google.com/machine-learning/glossary?authuser=0#L2_regularization

- Regularization

<https://developers.google.com/machine-learning/glossary?authuser=0#regularization>

[Help Center](https://support.google.com/machinelearningeducation?authuser=0) (<https://support.google.com/machinelearningeducation?authuser=0>)

[Previous](#)



[Video Lecture](#)

<https://developers.google.com/machine-learning/crash-course/regularization-for-simplicity/video-lecture?authuser=0>

[Next](#)

[Lambda](#)



<https://developers.google.com/machine-learning/crash-course/regularization-for-simplicity/lambda?authuser=0>

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies?authuser=0) (<https://developers.google.com/site-policies?authuser=0>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2020-02-10 UTC.